

Education

Master of Science (Financial Engineering)

[Stevens Institute of Technology](#)

School of Business (Division of Financial Engineering)

GPA (4.0 Scale): 3.8 | Awards: Provost's Master's Fellowship

Hoboken, NJ

May '19

Bachelor of Arts in Business Administration (Finance)

[University of Washington](#)

Michael G. Foster School of Business (Department of Finance and Business Economics)

Seattle, WA

Jun '17

Work Experience

Research Collaborator

[Rensselaer Polytechnic Institute](#)

SCIENCE Blockchain Project; Institute for Data Exploration and Applications (IDEA)

Troy, NY

Sep '21 - Present

- Assist in the research and development process of a trustworthy, accountable data sharing ecosystem for biomedical research.
- Lead efforts in developing cloud infrastructure to process and analyze data to build the underlying collaborator index. Currently exploring tokenomics policies for incentivizing data sharing in research networks.
- Help mentor undergraduate students by providing assistance with research organization, infrastructure development, project implementation, and publication drafting.

Software Engineer (Data Science)

[ExtraHop Networks](#)

Unusual Behaviors Group; Data Science R&D

Seattle, WA

Sep '19 - Present

- Designed and developed high throughput data pipelines to transport and load semi structured data across numerous globally distributed data centers. Experienced in creating enriched data sets to deliver insights to internal and external stakeholders.
- Produced numerous Data Science applications for both research and production use, with modern AWS serverless technologies and architectures. Applied modern Data Science algorithms and methodologies to Big Data to deliver insights to our customers.
- Formalized, engineered, and managed the core internal ETL pipeline for over 2 years, delivering value to numerous internal teams across the business. Scaled the core pipeline to handle an order of magnitude increase in data volume, while increasing functionality, due to increased adoption of ExtraHop's Cybersecurity product, Reveal(X) 360.
- Enhanced ExtraHop's core threat hunting ability as a detector writer, applying unsupervised learning algorithms to Big Data scale traffic flows in the cloud. Implemented numerous real-time detectors that act as early warning signals of potential bad actors on corporate networks across our customer base.
- Participated in recruitment efforts, and interviewed teammates before, and during the COVID pandemic. Assisted with mentoring and on-boarding new team members.

Research Assistant

[Stevens Institute of Technology](#)

Sensorimotor Control Laboratory; Department of Biomedical Engineering & Stevens Institute for Artificial Intelligence

Hoboken, NJ

Aug '18 - May '19

- Designed and implemented algorithms to assess and classify tremor severity in patients with late-stage Parkinson's Disease.
- Created a highly scalable and extensible web application to be used by the researchers in the lab during this project. This web application incorporated HIPAA-compliant data storage and access, as well as efficient cluster management with Docker and Kubernetes.

Summer Research Fellow

[RPI-IBM HEALS Research Center](#)

AI Horizons Network; IBM Research

Tetherless World Constellation; Rensselaer Polytechnic Institute

Troy, NY

May '18 - Aug '18

- Led the design and development of the PaperRank Framework, a methodology for deriving probabilistic community trust in academic publications. PaperRank utilized the PageRank algorithm, coupled with a Gamma Mixture Model applied to citation networks of academic publications. A proof-of-concept was implemented, from extraction to final trust score computation, analyzing over 14 Million articles from the NCBI PubMed Database.
- Formulated and implemented novel strategies for semantically-enhanced automated extraction of medical directives from Clinical Practice Guidelines (CPGs), for eventual inclusion in a knowledge graph of Diabetes diagnosis and treatment directives. Built the 'Guideline Explorer', a tool for efficiently visualizing and examining the American Diabetes Association's 2018 CPGs.
- Explored the field of 'Semantalytics', which lies at the intersection of Semantics and Analytics. Drafted a Vision statement for the future exploration of this novel field of research, through the lens of bioinformatics.

Publications

Learned Sectors: A fundamentals-driven sector reclassification project

2019

Rukmal Weerawarana, Yiyi Zhu, Yuzhen He

arXiv preprint; arXiv:1906.03935

Market sectors play a key role in enabling the efficient flow of capital through the modern Global economy. An analysis of existing sectorization heuristics show that they are not entirely quantitatively driven, but rather are highly subjective and rooted in dogma. To this end, we introduce a new fundamentals-driven Learned Sectors heuristic.

(Draft) Inferring Community Trust from Citation Graphs

2019

Jamie McCusker, Rukmal Weerawarana, Alexander New, Kristin P. Bennett, Deborah L. McGuinness

We introduce the PaperRank scoring algorithm; a proxy of scientific community trust in a given publication. This score is derived from the classic PageRank algorithm (applied to academic citation networks), in conjunction with a one-dimensional Gamma Mixture Model to normalize the PageRank scores on a 3-group publication notoriety heuristic.

Semantic Modeling of Cohort Descriptions in Research Studies

2018

Shruthi Chari, Rukmal Weerawarana, Oshani Seneviratne, Jamie McCusker, Deborah L. McGuinness, Amar Das

Knowledge Representation and Semantics Workshop; AMIA 2018 Annual Symposium

This research addresses a key challenge faced by physicians using Clinical Practice Guideline recommendations; determining how well idiosyncratic cohort evidence generalizes to the greater clinical population.

What is a Knowledge Graph?

2018

Jamie McCusker, John S. Erickson, Katherine Chastain, Sabbir Rashid, Rukmal Weerawarana, Marcello Bax, Deborah L. McGuinness

This work attempts to synthesize a clear and unambiguous definition of a 'Knowledge Graph' that conforms to current knowledge graph research, while constraining the research space that may be considered a knowledge graph.

Selected Projects

Precis | <https://precis.rukmal.me>

2021

Precis is an Ontology for modeling personal professional metadata. The extended Precis toolkit also includes a Pythonic search API for the Ontology, a JSON data loader, and an extensible templating engine.

fe621 | <https://git.rukmal.me/FE-621-Homework>

2019

fe621 is a Python library that provides functionality for lattice based derivative pricing models, exotic option pricing, Monte Carlo simulations, numerical differentiation and integration, and optimization.

reIndexer | <https://git.rukmal.me/reIndexer>

2019

reIndexer is a research tool for the backtest-driven evaluation of different sectorization heuristics, using a system of synthetic ETFs, and efficient portfolios of those synthetic ETFs.

HTKG and NYPD-Compstat-LD | <https://git.rukmal.me/NYPD-Compstat-LD>

2019

Collaborators: Ryan Hartman, Ayush Kalla, Kovid Shukla, Sanket Saharkar

HTKG (Human Trafficking Knowledge Graph), and NYPD-Compstat-LD (the main data ingestion engine) is a knowledge graph platform for linking suspected human trafficking advertisements with crime data from the NYPD to retroactively assess trends.

PaperRank Framework | <https://git.rukmal.me/PaperRank>

2018

The PaperRank framework is designed to enable bibliometrics and citation analysis of academic literature graphs. It is highly extensible, and designed to be corpus-agnostic; currently, it is configured for use with the NCBI PubMed database.

Selected Talks

Neural Ordinary Differential Equations

2020

<https://drive.google.com/file/d/1fqVH6GJe1TcRyD6tL2cDUXkkkhyRq4vY>

A literature review of Neural Ordinary Differential Equations by Chen et al., a new family of deep neural network models that parameterizes the hidden state of a neural network.

Knowledge Graph Fundamentals

2018

https://drive.google.com/file/d/1f21S_QZtm6aYiYLnXPA5opgsoBdIX7zX

An overview of the fundamental technology powering modern knowledge graphs, focusing on the concepts of semantic data, ontologies, and inference.

High Frequency Trading (HFT) - A Deep Dive

2017

<https://drive.google.com/file/d/1I7JuZhVzsAT84xe881LsbHbSIa-Ev7eF>

A deep dive into High Frequency Trading (HFT), covering market microstructure, exchange dynamics, regulatory implications, electronic order execution models, RegNMS, algorithmic trading, and popular HFT-driven strategies for exploiting arbitrage opportunities.